

Problem & User Research

Who DOCURA is for, what they are actually struggling with, what they do today instead, and what we must prove before we build.

DOCUMENT	Problem & User Research
PRODUCT	DOCURA — personal document intelligence layer
VERSION	1.0 — Draft for alignment
OWNER	Product Management
DATE	25 August 2026
PRECEDED BY	01 — Product Vision
FOLLOWED BY	03 — Product Principles & Scope

Evidence notice. No primary user research has been conducted for DOCURA at the time of writing. This document therefore separates what is **observed** from what is **assumed**, and every substantive claim carries an evidence tag. Nothing in this document should be cited as a validated finding until Section 15 has been executed. Its purpose is to make our current beliefs explicit, falsifiable, and prioritised for testing — not to imply research that has not happened.

– Contents

01	Purpose and Scope of this Document
02	Research Objectives
03	Method and Evidence Standard
04	The Problem Space
05	Root Causes
06	The Cost of the Problem
07	User Segments
08	Personas
09	Jobs to Be Done
10	Current-State User Journey
11	Pain Point Catalogue
12	User Needs
13	Trust, Privacy and Sensitivity
14	What Users Do Today Instead
15	Key Assumptions and Validation Plan
16	Risks Surfaced by the Research
17	Open Questions
18	Summary of Findings

01 Purpose and Scope of this Document

The Product Vision established what DOCURA intends to be. This document establishes **why it deserves to exist**: the specific problem, the specific people who have it, the shape that problem takes in their real behaviour, and the evidence standard we will hold ourselves to before committing engineering effort.

IN SCOPE

- Definition and decomposition of the core problem.
- User segmentation, personas, and jobs to be done.
- Current-state journey and the pain points along it.
- Functional, informational, emotional, and trust-related needs.
- Existing alternatives and workaround behaviour.
- Explicit assumptions, the risks they carry, and how we will test them.

OUT OF SCOPE

- Feature definition, prioritisation, and scope decisions — document 03.
- Interface design, interaction patterns, and visual system application.
- Technical architecture, model selection, and data infrastructure.
- Roadmap, delivery sequencing, and engineering tickets.
- Business model, pricing, and go-to-market strategy.

02 Research Objectives

This document exists to answer eight questions. Each is stated as a question rather than a conclusion, and each is answered later at the confidence level our current evidence actually supports.

TABLE 2.1 — RESEARCH OBJECTIVES AND WHERE THEY ARE ADDRESSED

#	Question	Addressed in
RO-1	What exactly is painful about filling official forms — is it the typing, the documents, the uncertainty, or the consequences?	\$4, \$11
RO-2	Why does the pain persist despite storage products, autofill, and government lockers already existing?	\$5, \$14
RO-3	Who feels this pain most acutely and most often?	\$7, \$8
RO-4	What is the user actually trying to accomplish, independent of any solution?	\$9
RO-5	Where in the journey does the user lose the most time, and where do they lose the most confidence? Are they the same place?	\$10, \$11

#	Question	Addressed in
RO-6	What would a user have to believe about a product before uploading their identity documents to it?	§13
RO-7	What are people doing today instead, and why is it good enough to have survived?	§14
RO-8	Which of our beliefs, if false, would invalidate the product — and how do we find out cheaply?	§15, §16

03 Method and Evidence Standard

Product documents fail most often not by being wrong, but by presenting belief in the grammar of fact. To prevent that, every substantive claim in this document carries one of three evidence tags.

TABLE 3.1 — EVIDENCE TAGS USED THROUGHOUT THIS DOCUMENT

Tag	Meaning	How it may be used
DESK	Directly observable without asking a user: form structures, portal behaviour, published requirements, product capabilities of existing tools.	May be relied on for design decisions. Must be re-checked when portals change.
ASM	A reasoned assumption about user behaviour, motivation, or priority. Plausible, unverified.	May guide exploration and prototypes. May not justify committed scope until tested.
VAL	Confirmed by primary research with real users, with sample size and method recorded.	May justify committed scope. No claim in version 1.0 of this document carries this tag.

3.1 Method used for this version

- **Structural analysis of the form domain** DESK — decomposing what application forms actually demand: field types, format constraints, document requirements, dropdown taxonomies, and file-upload rules.
- **Capability analysis of existing alternatives** DESK — what document lockers, cloud storage, browser autofill, and password managers can and cannot do, established from their documented behaviour.
- **Assumption mapping** ASM — making our beliefs about user behaviour explicit and ranking them by the damage they would cause if wrong.

3.2 Method planned before scope is committed

Section 15 specifies the primary research required to convert the assumptions in this document into findings. Until that research is executed, this document is a **hypothesis register**, and should be read as one.

STANDARD WE ARE HOLDING

No feature enters committed scope on the strength of an **ASM** claim alone. Either the assumption is tested first, or the feature is built as a deliberately reversible experiment and labelled as such.

4.1 Problem statement

CORE PROBLEM

People are repeatedly required to prove facts about themselves that they have already documented. The proof exists, is owned by them, and is final — but it is stored as images and paper rather than as information, so every new form forces a manual re-translation from document to field. The work is unskilled, unavoidable, high-volume, and disproportionately punished when it goes wrong.

4.2 The problem has three distinct layers

Treating this as a single problem is the mistake most adjacent products make. It is three problems that happen to occur in sequence, and each one defeats a different existing solution.

TABLE 4.1 — DECOMPOSITION OF THE CORE PROBLEM

#	Layer	What goes wrong	Why existing tools fail here
L1	Retrieval DESK	The right document, in the right version, format, and file size, cannot be located quickly. Documents are spread across phone galleries, chat threads, several email accounts, cloud drives, institutional portals, and physical folders.	Storage tools solve location only if the user filed things correctly in the first place. They do not solve version ambiguity or format constraints.
L2	Translation DESK	The information inside the document must be read by a human and retyped into fields whose labels, order, and formats differ on every portal. The same fact is entered dozens of times a year.	Autofill matches field names against saved strings. It has no access to the contents of a document and no model of what a field means.
L3	Judgement ASM	Several points in a form require a decision, not data: which category applies, which of two certificates is the right one, whether a declaration is true, whether to submit. These cannot be answered by lookup.	Autonomous agents optimise for completion and therefore guess. On applications with deadlines and legal declarations, a confident guess is worse than a question.

DOCURA's thesis is that a product must solve all three layers to be useful, and must solve the third one **differently in kind** — by refusing to automate it.

4.3 Why the problem is invisible

This work is rarely counted as work. It happens at night, before deadlines, alone, and produces nothing anyone sees. There is no line item for it, no one is measured on it, and the person doing it usually blames themselves for being disorganised rather than blaming the system for asking the same question forty times. ASM That invisibility is why

the problem has persisted at scale without a dedicated product — and it is also a marketing risk we address in Section 16.

05 Root Causes

Fixing symptoms produces a better filing cabinet. Fixing causes produces DOCURA. Five causes explain the persistence of the problem.

TABLE 5.1 — ROOT CAUSES AND THEIR IMPLICATIONS

#	Root cause	Explanation	Implication for DOCURA
RC-1	Documents are images, not data DESK	A scanned marksheet is a picture of information. Nothing in the file is queryable, comparable, or transferable without a human reading it.	Comprehension — not storage — is the foundational capability. Everything else depends on it.
RC-2	No shared vocabulary between forms DESK	The same fact appears as "Father's Name", "Guardian Name", "Parent/Guardian", and "Name of Father (as per 10th certificate)" across four portals, with different validation rules.	Field matching must be semantic, not string-based. A label dictionary is insufficient.
RC-3	Format constraints are arbitrary and unstated DESK	File type, dimensions, and size limits vary by portal and are often discovered only on rejection. Photographs and signatures are the most common casualties.	Document preparation is part of the problem, not a separate utility. Storing the original is not enough.
RC-4	Consequences are asymmetric ASM	Completing a form correctly earns nothing; completing it incorrectly can cost an admission, a scholarship, or a year. Users therefore over-check rather than move fast.	Speed alone is not the value. Verifiability and reversibility are what reduce the real cost.
RC-5	Trust prevents delegation ASM	The documents involved are the most sensitive a person owns. Any tool asking for all of them at once faces a trust barrier before it can deliver any value.	Trust must be earned incrementally, and the product must be valuable before it is fully populated.

06 The Cost of the Problem

We deliberately do not quote time or money figures here. Precise-sounding numbers without primary measurement are the most common way a research document becomes misleading. Instead we name the **cost categories** and specify how each will be measured during validation.

TABLE 6.1 — COST CATEGORIES AND THEIR MEASUREMENT PLAN

Cost category	How it manifests	How we will measure it
---------------	------------------	------------------------

Cost category	How it manifests	How we will measure it
Time	Hours per application spent locating, preparing, retyping, and re-uploading. Repeated across many applications per season.	Timed observation of a real form completion, before and after DOCURA (\$15, S-2 and S-5).
Rework	Applications corrected, resubmitted, or restarted because of a format, mismatch, or upload error.	Self-reported incidence in interviews; error counts in observed sessions.
Opportunity	Applications not made at all — scholarships skipped, deadlines missed — because the effort exceeded the perceived odds.	Interview probe on abandoned applications in the last twelve months.
Cognitive and emotional	Deadline anxiety, low-grade dread of official processes, and self-blame for disorganisation.	Qualitative coding of interview language; confidence rating before and after submission.
Dependency	Reliance on a parent, sibling, senior, or cyber café to complete one's own application.	Interview probe on who else was involved in the last three applications.

WHY THIS MATTERS FOR POSITIONING

If validation shows the dominant cost is **time**, DOCURA is an efficiency product and competes on speed. If it shows the dominant cost is **rework and anxiety**, DOCURA is a confidence product and competes on correctness. These lead to materially different products. Resolving this is the single highest-value outcome of the first research round.

07 User Segments

Segments are ordered by **frequency × severity × reachability**, not by market size. A large segment that files two forms a year is a worse starting point than a small segment that files twenty.

TABLE 7.1 – SEGMENT PRIORITISATION

Segment	Form frequency	Consequence of error	Document volume	Priority and rationale
Students and recent graduates	Very high — clustered in admission and scholarship seasons	Severe — a missed deadline can cost a year	High and growing each semester	Primary. Highest frequency, highest stakes, most documents, and a naturally recurring trigger calendar.
Early-career professionals	Moderate — job changes, onboarding, verification, registrations	High — delays affect employment and compliance	Stable, with periodic additions	Secondary. Same behaviour, lower frequency; reachable as students graduate into it.
Household document managers	Moderate to high — acting on behalf of several people	High — errors affect dependants	Very high — multiplied by family size	Tertiary. Strong pain, but multi-person records raise consent and access questions we should not solve first.
Institutions and bulk processors	Continuous	Operational	Very high	Explicitly excluded. A different trust model; serving them early would compromise the individual-first architecture.

7.1 Beachhead recommendation

Students in their final two years of undergraduate study ASM are the recommended beachhead. They are simultaneously applying for internships, higher studies, scholarships, and examinations; their document set is at its most complete and most frequently demanded; and they are the population most willing to adopt a new tool during a deadline crunch. This recommendation is an assumption and is listed for validation as A-3 in Section 15.

08 Personas

These personas are **design instruments, not research subjects** ASM. They are composite constructions used to keep decisions concrete and to give the validation interviews a hypothesis to break. They must be revised — or discarded — after the first research round.

PERSONA 1 • PRIMARY SEGMENT

The Applicant in Season

CONTEXT	Final-year undergraduate student. Applying to internships, postgraduate programmes, and at least two scholarships in the same three-month window.
DOCUMENT SET	Aadhaar, PAN, 10th and 12th marksheets, six semester results, student ID, photograph, signature, category and income certificates, resume, several course certificates.
CURRENT BEHAVIOUR	Keeps a WhatsApp self-chat and a phone folder of scans. Re-crops the photograph for each portal. Asks a parent for details that are not memorised.
PRIMARY GOAL	Submit every application before its deadline without an error that disqualifies them.
BIGGEST FEAR	Discovering after the deadline that something was wrong, mismatched, or never uploaded.
SUCCESS LOOKS LIKE	Finishing an application in one sitting and feeling certain it is correct.

"I know I have it somewhere — I just don't know which version they want."

Illustrative phrasing of the hypothesised mindset, not a research quotation.

PERSONA 2 • SECONDARY SEGMENT

The Early-Career Professional

CONTEXT	Two to five years into work. Changing jobs, completing employer onboarding and background verification, and handling periodic statutory and financial registrations.
DOCUMENT SET	Identity and address proof, educational records, previous employment and relieving letters, payslips, bank and tax records, professional certifications.
CURRENT BEHAVIOUR	A cloud drive folder that is organised in intention and chaotic in practice. Emails documents to themselves as a retrieval mechanism.
PRIMARY GOAL	Clear document requests quickly so they do not block onboarding, payroll, or a start date.
BIGGEST FEAR	Looking unprofessional or delaying a process because they could not produce a document on request.
SUCCESS LOOKS LIKE	Answering any document request the same day, from a phone, without searching.

The Household Document Manager

CONTEXT	The family member who handles paperwork for everyone — a parent's pension or insurance, a sibling's admission, a child's school forms.
DOCUMENT SET	Multiple people's identity, education, property, medical-adjacent administrative, and financial records.
CURRENT BEHAVIOUR	A physical folder plus a partially mirrored digital copy. Deep institutional knowledge held entirely in their head.
PRIMARY GOAL	Never be the reason a family member's application failed.
BIGGEST FEAR	Losing an original, or being unreachable when a document is urgently needed.
SUCCESS LOOKS LIKE	Anyone in the family getting what they need without waiting for them.

Note: this persona introduces multi-person records, delegated access, and third-party consent. These are real needs but they are **deliberately deferred** — solving them early would force consent architecture decisions before the single-user model is proven.

09 Jobs to Be Done

Jobs are stated independently of DOCURA so that they remain valid even if our solution changes. They are ordered by how central they are to the product thesis.

JOB 1 — FUNCTIONAL, PRIMARY

When I am filling out an official form, I want the facts and documents I have already collected to be applied to it correctly, so I can complete the application without translating my own records by hand.

JOB 2 — FUNCTIONAL

When a form asks for a document in a specific format, I want the correct document produced in the correct format, so a technical constraint does not become a reason my application fails.

JOB 3 — INFORMATIONAL

When I need a fact about myself that lives inside a document, I want to find it in seconds without opening files one by one, so a small question does not become a ten-minute search.

JOB 4 — EMOTIONAL

When I submit something with a deadline and consequences, I want to feel certain that it is correct and complete, so I am not carrying anxiety after I have already submitted.

JOB 5 — CONTROL

When something is being done on my behalf, I want to remain the one who decides what is disclosed and what is agreed to, so speed never costs me authorship of my own declarations.

JOB 6 — PREPARATORY

When a season of applications is approaching, I want my records to already be complete and current, so preparation is not something I do under deadline pressure.

THE DISTINGUISHING JOB

Job 5 is the one no competing product treats as a job at all. Storage products have no concept of acting on the user's behalf; autofill and autonomous agents treat consent as friction to be minimised. Treating control as a job to be served — rather than an obstacle to be reduced — is what makes DOCURA a different product rather than a better one.

10 Current-State User Journey

The journey below traces a single application from trigger to aftermath. It is constructed from the structural analysis of forms DESK combined with assumed user behaviour ASM, and is the primary artefact to be corrected by the diary study in Section 15.

TABLE 10.1 — CURRENT-STATE JOURNEY, SINGLE APPLICATION

#	Stage	What the user does today	Where it hurts	Opportunity for DOCURA
S1	Trigger	Learns of a deadline from a notice board, a group chat, a senior, or a portal announcement.	Often late. Preparation time is compressed from the start.	Low — outside our control initially. Revisit once the record is populated.
S2	Requirement discovery	Opens the form or notification and works out which documents and details are required.	Requirements are scattered across the form, a PDF notification, and validation errors found only on submission.	High. Reading the form's full requirements up front is a distinct, valuable act.
S3	Document gathering	Searches phone gallery, chat threads, email, drive, and physical folders. Sometimes re-downloads from an institutional portal.	Slowest and most frustrating stage. Version ambiguity: which scan, which attestation, which year.	Very high. This is the retrieval layer (L1) and the clearest source of felt relief.
S4	Document preparation	Crops, rotates, converts, compresses, and renames files to meet stated limits.	Fiddly, tool-dependent, and frequently repeated per portal. Photograph and signature are worst.	Very high. Mechanical, rule-based, and fully automatable with certainty.
S5	Data entry	Reads values off documents and types them into fields; resolves dropdowns by scanning long option lists.	Repetitive and error-prone. Dropdown taxonomies differ per portal and rarely match the document's wording.	Very high. This is the translation layer (L2) and the core of the extension.
S6	Ambiguity resolution	Pauses on fields they are unsure about. Asks a parent, a senior, or a group chat. Sometimes guesses.	Breaks flow entirely. Guessing here is where disqualifying errors are introduced.	High. The judgement layer (L3): surface the choice with context rather than resolving it silently.
S7	Declarations and consent	Ticks agreement boxes and declarations, usually without reading them.	Legal exposure accepted without comprehension — a real risk, not merely a formality.	Deliberately unautomated. Opportunity is to make the moment clearer, never faster.
S8	Review and submit	Re-reads the filled form, often several times, before pressing submit.	Anxious, slow, and unreliable — the user is proofreading their own typing.	High. A verifiable review, showing each value's source document, replaces anxious re-reading.
S9	Aftermath	Saves a confirmation screenshot or PDF; hunts for it again weeks later when a hall ticket or status update is needed.	Confirmations become new unfiled documents, restarting the original problem.	Moderate. Closing this loop turns a one-time tool into a permanent record.

JOURNEY INSIGHT

Time is lost mainly at **S3, S4 and S5**. Confidence is lost mainly at **S6 and S8**. They are not the same stages — which means a product that only saves time will not resolve the anxiety, and a product that only reassures will feel slow. DOCURA must address both, and should measure them separately (RO-5).

11 Pain Point Catalogue

Pains are scored on frequency and severity to give a defensible ordering for scope discussions in document 03. Scores are **current best judgement** ASM and are to be replaced with measured values after validation.

TABLE 11.1 – PAIN POINTS, ORDERED BY COMBINED PRIORITY (H = HIGH, M = MEDIUM, L = LOW)

ID	Pain point	Freq.	Sev.	Journey stage	Note
P-01	Cannot find the right version of a document quickly	H	H	S3	Compounded by unclear file names and duplicates across devices.
P-02	Re-entering the same personal details on every portal	H	M	S5	Highest-volume pain; individually small, cumulatively dominant.
P-03	File format and size rejections on upload	H	H	S4	Often discovered at submission, under time pressure.
P-04	Dropdown options that do not match the document's wording	H	H	S5, S6	A frequent source of silent, disqualifying mismatches.
P-05	Uncertainty about which of several similar documents is required	M	H	S3, S6	Provisional vs final, attested vs plain, consolidated vs per-semester.
P-06	Name and detail mismatches across documents	M	H	S5, S8	Initials, expansions, spellings. Can invalidate an otherwise correct application.
P-07	Anxiety that something was missed after submitting	H	M	S8, S9	Persists after the task ends; rarely resolvable by the user.
P-08	Session timeouts and lost progress on slow portals	M	H	S5	Forces a full re-entry; strongly increases abandonment.
P-09	Needing another person to supply a detail	M	M	S6	Creates dependency and delay, often overnight.
P-10	Signing declarations without reading them	H	M	S7	Low felt pain, real exposure. A responsibility for us, not a demand from users.
P-11	Confirmations and receipts becoming lost documents	M	M	S9	Regenerates the original problem after every application.
P-12	Preparing the same document differently for each portal	M	M	S4	Repeated mechanical work with no learning benefit.

P-01 through P-04 are high on both axes and cluster in S3–S5. They constitute the **minimum credible problem** DOCURA must solve to be worth adopting. P-10 is the outlier: users do not experience it as pain, but it is where automation would do the most harm — it is in scope because of our philosophy, not because of demand.

12 User Needs

Needs are grouped by type, because each type is satisfied by a different kind of design decision and validated by a different kind of evidence.

TABLE 12.1 — NEEDS BY CATEGORY

ID	Category	Need	Consequence if unmet
N-01	Functional	Correct values placed in correct fields, in the format the portal expects.	The product is a novelty; users revert to typing.
N-02	Functional	The right document attached in the right slot, prepared to the stated constraints.	Half the work remains manual; adoption stalls after the first form.
N-03	Informational	Fast retrieval of any fact or document from one's own record.	The vault is a place things go to be forgotten.
N-04	Informational	Visibility of what a form requires before starting it.	Users still discover requirements through failure.
N-05	Verification	Ability to see where each filled value came from.	Users re-check everything manually; the time saving evaporates.
N-06	Emotional	Justified confidence at the moment of submission.	Anxiety persists; the product solves speed but not the felt problem.
N-07	Control	Final authority over sensitive disclosures, declarations, and submission.	Trust collapses on the first surprising action, and does not return.
N-08	Trust	Confidence that the record is private, inspectable, correctable, exportable, and deletable.	Users never populate the vault, so nothing downstream can work.
N-09	Maintenance	Low-effort ways to keep the record current as new documents arrive.	The record decays; accuracy silently degrades over time.
N-10	Continuity	Access from wherever the form is being filled, including a shared or borrowed device.	The product fails at the exact moment of need.

N-05 and N-08 are the load-bearing needs. Without traceability, automation creates work instead of removing it, because a careful user will re-verify everything anyway. Without trust, the vault is never populated, and every other capability is unreachable. Both must be designed for first, not added later. ASM

13 Trust, Privacy and Sensitivity

DOCURA asks for the most sensitive documents a person owns, before it has delivered any value. This is the product's hardest adoption problem, and it is a research problem before it is a security problem.

13.1 The trust sequence

Users do not evaluate trust once; they evaluate it repeatedly, and the order matters. ASM

- **Before uploading anything** — Who can see this? Where does it go? What happens if I stop using this?
- **On first automated action** — Did it do exactly what I expected, and can I see why it did it?
- **On the first sensitive field** — Did it ask me, or did it decide for me?
- **On the first mistake** — Was the mistake visible and correctable, or did I find it later, elsewhere?

Trust is granted incrementally and withdrawn instantly. A single unexpected autofill of a sensitive field can end adoption permanently, regardless of how correct every prior action was.

13.2 Sensitivity is not uniform

Treating all stored information as equally sensitive produces either excessive friction or insufficient protection. A working classification for research purposes: ASM

TABLE 13.1 — WORKING SENSITIVITY CLASSIFICATION

Tier	Examples	Expected user expectation
Routine	Name, date of birth, education history, marks, addresses already public on the application.	Fill it silently. Asking each time is experienced as friction, not care.
Sensitive	Government identifier numbers, bank details, signature images, category and income details.	Show it, confirm it, then fill it. Never place it without an explicit action.
Consequential	Legal declarations, agreements, affidavit-style statements, final submission.	Never automated under any confidence level. The user performs the act themselves.

The boundary between *routine* and *sensitive* is the most important unresolved question in this document, because it determines how often DOCURA interrupts. Interrupt too often and the product is slower than typing; interrupt too rarely and it violates the principle it was built on. This boundary must be set by users, not by us — it is assumption A-5 in Section 15.

RESEARCH IMPLICATION

Trust research must be conducted with users' **real documents**, not test data. Willingness to upload a sample file tells us nothing about willingness to upload an actual identity document. Any study that substitutes dummy data will produce a false positive on our riskiest assumption.

14 What Users Do Today Instead

Every workaround that survives is doing a real job. Understanding what each one gets right explains both why the problem persists and what DOCURA must not lose.

TABLE 14.1 – CURRENT ALTERNATIVES AND THEIR LIMITS

Alternative	What it does well	Where it fails	What DOCURA must preserve
Phone gallery and screenshots	Instant capture; always in hand; zero setup.	No structure, no search, no versioning, no format control.	Capture must be as fast as taking a photograph.
Messaging self-chat	Effortless, chronological, searchable by memory of when it was sent.	Buried in unrelated messages; compression degrades documents; no requirement awareness.	Adding a document must require no filing decision.
Cloud drive folders	Durable, cross-device, shareable.	Organisation is manual and decays; no comprehension of contents.	Durability and access from any device.
Government document lockers	Authoritative, issuer-verified copies.	Limited to issued documents; no form awareness; no personal additions such as resumes or photographs.	Respect for authoritative sources; never replace them, complement them.
Browser and password-manager autofill	Genuinely fast on recognised, simple fields.	String matching only; no documents; no dropdown reasoning; no confidence model.	The speed of the common case must not regress.
Cyber cafés and paid form-fillers	Transfers the whole problem to someone experienced; handles printing and scanning.	Costly, requires sharing all documents with a stranger, and creates dependency.	The feeling of handing the work to someone competent.
A knowledgeable family member	Knows the history, catches mismatches, is trusted absolutely.	Not always available; single point of failure; creates dependency.	Judgement offered, not imposed — the behaviour of a trusted advisor.

COMPETITIVE INSIGHT

Our true competitor is not another product; it is **the self-chat plus a family member**. That combination is free, trusted, and already installed. DOCURA must beat it on retrieval speed and correctness without ever asking for more trust than it has earned.

15 Key Assumptions and Validation Plan

These are the beliefs this document rests on, ordered by **damage if wrong**. Each has a stated test and a falsification criterion, so the research can disconfirm rather than merely confirm.

TABLE 15.1 – ASSUMPTION REGISTER

ID	Assumption	Damage if wrong	How we test it	Falsified if
A-1	Users will upload real identity documents to a new product from an unknown team.	Fatal	Trust interviews plus an observed first-upload session using their own documents.	A majority decline to upload identity documents, or upload only low-stakes files.
A-2	Being asked at ambiguous points is experienced as trustworthy, not as the product failing.	Fatal	Prototype comparison: a version that asks versus one that silently decides.	Users prefer silent completion even after being shown an induced error.
A-3	Final-year students are the right beachhead.	High	Frequency and severity comparison across the three segments in interviews.	Another segment shows higher frequency and equal willingness to adopt.
A-4	Document retrieval and preparation (S3–S4) hurt more than data entry (S5).	High	Timed diary study across at least three real applications.	Data entry dominates measured time and reported frustration.
A-5	Users agree with our routine / sensitive / consequential boundary.	High	Card sort of real field types by required confirmation level.	Users' boundary differs materially, or varies too widely to generalise.
A-6	One-time setup effort is acceptable in exchange for later speed.	Medium	Observed onboarding with a real document set; measure drop-off point.	Users abandon before completing a minimum viable record.
A-7	Document comprehension is accurate enough on real-world scans to be trusted.	Medium	Technical evaluation against a corpus of genuinely messy user documents.	Accuracy on real scans is low enough that review costs more than typing.
A-8	A browser extension is an acceptable and reachable delivery mechanism.	Medium	Observe where applications are actually completed — desktop, mobile, or shared device.	A majority complete applications on mobile or on devices where extensions cannot be installed.

15.1 Proposed research programme

TABLE 15.2 – STUDIES, SEQUENCED

#	Study	Participants	Answers	Output
---	-------	--------------	---------	--------

#	Study	Participants	Answers	Output
S-1	Trust and willingness interviews	12–15	A-1, A-5, RO-6	Trust sequence map; validated sensitivity tiers.
S-2	Diary study across a real application season	8–10	A-4, RO-1, RO-5	Corrected journey with measured time per stage.
S-3	Form teardown of high-traffic portals	n/a — desk	RO-2, A-7, A-8	Field taxonomy, format-constraint catalogue, feasibility read.
S-4	Card sort of field sensitivity	15–20	A-5	The interruption boundary, set by users.
S-5	Prototype test: asks versus decides	10–12	A-2, A-6	Evidence for or against the core philosophy.
S-6	Comprehension evaluation on real scans	n/a — technical	A-7	Accuracy baseline and confidence-threshold guidance.

Sequencing note. S-1 and S-5 must run before scope is committed, because A-1 and A-2 are the only two assumptions whose failure invalidates the product rather than merely redirecting it. S-3 and S-6 can run in parallel, since they require no participants.

16 Risks Surfaced by the Research

TABLE 16.1 — RISKS AND MITIGATIONS

ID	Risk	Why it is serious	Mitigation direction
R-1	Cold-start trust barrier	The product requires the most sensitive documents before it has proven any value. Adoption may fail before the product is ever experienced.	Deliver value at low document counts; let the user start with a single low-stakes document.
R-2	Setup cost exceeds perceived benefit	Populating a record is front-loaded work; the payoff arrives later. Users abandon during onboarding.	Make the first form completion possible with a minimum viable record; capture the rest opportunistically.
R-3	Interruption fatigue	If the sensitivity boundary is set too cautiously, asking becomes nagging and the product feels slower than typing.	Set the boundary from S-4, and let users adjust it within limits that never include consent.
R-4	Silent comprehension errors	A misread value filled confidently is worse than no automation, because it defeats the user's own proofreading.	Traceability by default (N-05); confidence thresholds tuned to prefer asking over acting.

ID	Risk	Why it is serious	Mitigation direction
R-5	Portal volatility	Forms change without notice; brittle matching degrades silently and erodes trust exactly when stakes are highest.	Semantic matching over structural; degrade to asking rather than to guessing.
R-6	Invisible problem, invisible product	Users blame their own disorganisation rather than the system, so they may not search for a solution at all.	Reach users at the moment of pain — deadline seasons — rather than through category demand.
R-7	Mobile-first reality	If most applications are completed on phones, a browser extension reaches the wrong surface.	Resolve early via A-8; treat the delivery surface as an open question, not a settled decision.

17 Open Questions

- **Q-1.** Where do users actually complete applications — personal laptop, shared desktop, mobile, or a cyber café? This determines the product's primary surface and is currently unanswered.
- **Q-2.** How often does a user encounter a genuinely new form versus a repeat of one they have filled before? Repetition rate determines how much value comes from learning.
- **Q-3.** Is the felt problem seasonal or continuous? Seasonality changes retention expectations and what "active use" should mean.
- **Q-4.** Do users want DOCURA to detect inconsistencies between their own documents — for example a name mismatch — or would that feel like being audited?
- **Q-5.** What is the minimum record size at which the product first feels useful? This sets the onboarding target.
- **Q-6.** Would users accept a fully local-only mode with reduced capability in exchange for stronger privacy guarantees?
- **Q-7.** How do users currently decide which of two similar documents to submit — and what would they need to see to make that decision faster?

18 Summary of Findings

WHAT WE BELIEVE WITH REASONABLE CONFIDENCE DESK

- The problem is real, structural, and three-layered: retrieval, translation, and judgement. Existing categories each solve one layer and none solve all three.
- Documents are images rather than data, which is the root cause that makes every downstream problem persist.
- Format and dropdown constraints are a significant, fully mechanical source of failure that no current tool addresses.
- Consent and declaration steps are structurally unsuited to automation, regardless of technical capability.

WHAT WE BELIEVE BUT HAVE NOT TESTED ASM

- Retrieval and preparation hurt more than typing does.
- Students in their final years are the correct beachhead.
- Being asked at ambiguous moments increases trust rather than signalling failure.
- Users will upload real identity documents to a new product on the strength of stated guarantees.

WHAT WOULD CHANGE THE PRODUCT ENTIRELY

- If A-1 fails, the vault model is wrong and the product must be re-conceived around per-session, non-retained documents.
- If A-2 fails, the core philosophy is a preference of ours rather than a user need, and the positioning must be rebuilt.
- If A-8 fails, the browser extension is the wrong delivery surface and the mobile experience becomes the primary product.

RECOMMENDATION

Proceed to product principles and scope definition, but **hold feature commitment** until studies S-1 and S-5 have reported. Those two studies test the only assumptions whose failure would invalidate DOCURA rather than redirect it, and both can be run with a small number of participants at low cost. Everything else in this document is safe to design against provisionally, provided each item is carried forward with its evidence tag intact.

END OF DOCUMENT · DOCURA 02 — PROBLEM & USER RESEARCH · V1.0